

Physically Interpretable World Models for Partially Observable Driving: Encoding and Predicting the Road Environment under Weak Supervision

[AUTHOR LIST — TO BE COMPLETED BEFORE SUBMISSION]^a

^a*Trustworthy Engineered Autonomy (TEA) Lab, University of
Florida, Gainesville, Florida, USA*

Abstract

Learning predictive models from high-dimensional observations is fundamental for safe autonomous cyber-physical systems, yet the latent states of standard world models lack physical interpretability. Our conference work introduced the Physically Interpretable World Model (PIWM), which aligns latent representations with physical quantities and constrains their evolution through partially known dynamics under distribution-based weak supervision. That formulation assumes the vehicle state is a complete Markov description of the task, which holds for fully observable benchmarks such as CartPole and Lunar Lander but fails for driving: two vehicles with identical position, heading and speed follow divergent futures when the road ahead differs, and a forward-facing camera reveals that road only partially. This article extends PIWM to partially observable driving along three axes. First, we move to two partially observable driving case studies: a simulated top-down *CarRacing* track and a *real, physical DonkeyCar*. Second, we introduce a physically interpretable *encoding* of the road environment, factorizing the interpretable latent into a vehicle-relative component and a road-context component—lateral deviation, vehicle–road heading error and curvature—inferred from onboard images. Third, we introduce a physically interpretable *prediction* of that environment: a split-dynamics model that advances the vehicle through structured kinematics while forecasting how the road context evolves. On the physical DonkeyCar, where road context is supervised only during training and must be recovered from the front camera at test time, PIWM attains 0.463 m position error over a 100-step rollout against 0.703 m for the strongest baseline under five-fold cross-validation, ahead in every fold, with less than

half the fold-to-fold variation of any baseline—and with every state variable physically meaningful.

Keywords: physically interpretable world models, partial observability, autonomous driving, road-context encoding, weak supervision, trajectory prediction

1. Introduction

Accurate and robust trajectory prediction from high-dimensional sensor data is a foundational challenge for autonomous cyber-physical systems (CPS). Modern *world models* [1] address this challenge by encoding observations into compact latent representations and predicting their temporal evolution, a paradigm that has proven effective for long-horizon prediction, planning, and control [2, 3, 4]. Despite strong empirical performance, the latent states of these models typically lack a clear connection to the underlying physical state of the system. In safety-critical settings such as autonomous driving and robotics, this opacity limits trustworthiness: without interpretable states, it is difficult to detect unsafe future configurations, diagnose prediction failures, or apply formal safety-monitoring techniques [5, 6, 7].

Our prior conference work introduced the Physically Interpretable World Model (PIWM) [8], a framework that aligns latent representations with physical quantities and constrains their temporal evolution through a partially known dynamics model. PIWM defines physical interpretability through two complementary properties: the latent state must correspond to meaningful physical variables (state semantics), and its temporal evolution must follow physically consistent dynamics (temporal correctness). A key feature of PIWM is *distribution-based weak supervision*: rather than requiring precise ground-truth state labels, the method accepts state estimates in the form of sample distributions, as naturally produced by real-world sensing pipelines such as GPS or radar. Experiments on CartPole, Lunar Lander, and DonkeyCar demonstrated that PIWM recovers accurate physical parameters and achieves long-horizon prediction accuracy well beyond data-driven baselines. A systematic comparison of intrinsic versus extrinsic architectures and continuous versus discrete latent spaces showed that an extrinsic architecture with a quantized (VQ-based) latent yields the most robust and physically grounded representations.

The challenge: driving is partially observable

Applying PIWM to realistic driving exposes a fundamental limitation of the original formulation, and motivates the extensions in this article. On CartPole and Lunar Lander, the system state completely specifies what will happen next: the dynamics are fully observable and Markov, so a world model that tracks the vehicle (or pole, or lander) state captures everything relevant to the future. Driving is different. Consider two vehicles with identical position, heading, and speed. If one is approaching a tight curve while the other is on a straight segment, their future trajectories diverge even under the same control inputs. The vehicle’s physical state $x_t = (x, y, \psi, v)$ does not capture this difference. The missing information is the *road context*: the shape of the lane ahead (its curvature) and the vehicle’s lateral and angular relationship to it, none of which is part of any standard vehicle-state representation. As a result, the vehicle state alone is *non-Markov* with respect to future motion: two identical states lead to different futures depending on the unobserved road geometry. Applying the original PIWM formulation to such tasks implicitly treats the system as fully observable when it is not, causing the dynamics model to learn an underdetermined, ill-posed mapping.

This challenge has two distinct facets, and overcoming each requires a separate mechanism. *(i) The road environment must be encoded.* The relevant road context is not given as a clean state vector; it is latent in the pixels and, for a forward-facing camera, only partially visible at any instant. A physically interpretable driving world model must therefore *infer* the geometry of the surrounding road from raw images, and represent it in physically meaningful terms rather than as opaque features. *(ii) The road environment must be predicted.* Long-horizon rollout requires anticipating how the road context itself evolves as the vehicle advances: the lane that is far ahead now becomes the lane the vehicle is on a few steps later. A world model that predicts only the vehicle state, treating the road as static or as noise, accumulates error rapidly because it cannot anticipate the geometry it is about to enter. Predicting the road, like encoding it, should be physically interpretable so that the forecasted environment can be inspected, trusted, and monitored.

These two facets, encoding and predicting the external environment, are what distinguishes a driving world model from the vehicle-centric formulation of the conference paper, which modeled only the state of the controlled body. They are also why partial observability cannot be papered over by simply adding capacity to a black-box predictor: the missing variable is structured,

physical, and causally upstream of the vehicle’s future, so it deserves an explicit, physically grounded treatment.

Our approach

We resolve partial observability by reformulating the interpretable state as a structured composition of two complementary components. The *vehicle-relative state* z^v captures the ego vehicle’s displacement and heading change expressed in a local reference frame anchored at the beginning of each prediction window; this representation removes dependence on global coordinates and improves transferability across trajectories. The *road-context state* z^r captures task-relevant environmental structure: the cross-track error (lateral deviation from the lane center), the vehicle–road heading error, and the local road curvature. Together, (z^v, z^r) form a Markov description of the driving task, restoring the validity of the predictive world-model framework. The vehicle-relative component is governed by a structured kinematic model whose unknown actuation parameters are recovered during training; the road-context component is governed by a geometry-driven transition that advects the road past the vehicle as it advances, rather than treating the road as a free latent that can drift. This split-dynamics formulation provides appropriate inductive biases for each component and decouples the faster-varying vehicle motion from the slower-varying road context.

We validate the extended formulation in two new partially observable driving environments. In **CarRacing**, a top-down simulated racing environment, the road geometry is present in the image but absent from the vehicle state; the model must learn to infer and exploit road-context information from raw pixels. On the physical **DonkeyCar** platform, a real RC-scale car observes its environment through a front-facing camera and drives on a laboratory track. During training, an external localization system provides ground-truth road context as a privileged signal; at test time this signal is withheld, and the model must infer and roll out road context solely from the onboard front-facing image stream. This setup tests whether the factorized representation learned under privileged supervision transfers to the more challenging unaided real-world deployment setting. On the physical platform, over a 100-step (≈ 4.5 s) rollout and under five-fold cross-validation, PIWM with factorized state and split dynamics reaches 0.463 m position error against 0.703 m for the strongest data-driven or physics-informed baseline, ahead in every fold, and the gap widens with the horizon: the baselines diverge while PIWM stays bounded.

Contributions

This journal article extends the conference version [8] along three axes, which constitute its three contributions.

1. **Two partially observable driving case studies.** We move PIWM beyond the fully observable benchmarks of the conference paper to two new driving environments that are genuinely partially observable: a simulated top-down *CarRacing* track and a *real, physical DonkeyCar* platform driven on a laboratory track under onboard-camera observation. We formally characterize partial observability in driving, showing that the vehicle state is non-Markov with respect to future trajectory, and use these two case studies to demonstrate the practical consequences and the benefit of the proposed treatment.
2. **Physically interpretable encoding of the road environment.** Unlike the conference formulation, which modeled only the controlled body’s state, we model the *external environment*. We introduce a factorization of the interpretable latent into a vehicle-relative component z^v and a road-context component z^r , where z^r comprises physically meaningful descriptors of the road (cross-track error, vehicle–road heading error, and curvature). We show how to infer z^r from images using temporal context, including the forward-camera case where the road ahead is only partially visible, and how privileged supervision available only at training time can be used to learn this encoding.
3. **Physically interpretable prediction of the road environment.** We introduce a split-dynamics model that not only rolls out the vehicle with a structured kinematic model, but also *forecasts the road context* with a geometry-driven transition in which the road is advected past the vehicle by its own predicted motion. The predicted road context is itself physically interpretable, so the forecasted environment, not just the forecasted vehicle, can be inspected and monitored. We show that this geometric treatment, rather than propagating the road as a free latent, is what keeps long-horizon rollout bounded, and that the road context can be perceived from the onboard camera at a cost of roughly 2 cm in 100-step position error relative to a privileged map lookup.

The remainder of this article is organized as follows. Section 2 surveys related work. Section 3 reviews background on autonomous CPS, world models, and physical interpretability. Section 4 formalizes the partial observability

problem in driving and introduces the factorized state. Section 5 presents the PIWM architecture, the road-context encoder (Section 5.2), and the road-context predictor (Section 5.3). Section 6 describes the experimental setup, and Section 7 reports results on the benchmarks and on the real DonkeyCar. Section 8 discusses implications and limitations, and Section 9 concludes.

2. Related Work

We organize the discussion into five threads relevant to physically interpretable world modeling under partial observability: trajectory prediction and safe forecasting (Section 2.1), representation learning for dynamical systems (Section 2.2), world models (Section 2.3), physics-informed and interpretable learning (Section 2.4), and partial observability, privileged learning, and road structure (Section 2.5). We close each thread by positioning PIWM against it.

2.1. Trajectory Prediction and Safe Forecasting

Predicting future motion is central to safe planning and control [9]. Classical predictors propagate a low-dimensional state with hand-specified models: Kalman-filter formulations over a bicycle model of vehicle motion [10], and Hamilton–Jacobi reachability for sets of reachable states with formal guarantees [11, 12, 13]. These methods are transparent and certifiable, but they require an explicit dynamical model and a low-dimensional state, and they scale poorly to high-dimensional camera observations.

Learning-based predictors instead forecast directly from data. Social-pooling and generative sequence models capture multi-agent interaction [14, 15, 16], while graph- and attention-based models such as Trajectron++ achieve strong forecasting by conditioning on scene graphs and high-definition maps [17]. Hybrid schemes embed classical filters inside learned models, as in learned Kalman filtering [18]. These models, however, depend on handcrafted scene representations or HD-map annotations that are unavailable in many settings [19, 20], and their internal states carry no physical meaning.

A parallel line equips predictors with reliability guarantees without altering the predictor itself: conformal prediction and calibrated forecasting bound error or coverage [21, 22, 23, 24]; statistical model checking and predictive monitoring estimate the probability of property violations [25, 26, 27, 28]; and runtime and misbehavior monitors flag unsafe behavior online [29, 30, 31, 32]. These techniques are complementary to ours: they wrap a predictor

with guarantees but do not make its latent state physically meaningful. Compared with this body of work, PIWM learns interpretable representations directly from raw images under weak supervision, without handcrafted scene representations or test-time state labels, so that the forecast itself, both the vehicle and the road around it, can be inspected.

2.2. Representation Learning for Dynamical Systems

Compact latent representations underpin learning from high-dimensional observations, from classical autoencoders [33] to variational autoencoders (VAEs) [34], but standard objectives impose no physically meaningful structure. A large literature seeks more structured latents by encouraging disentanglement [35, 36, 37], causal structure [38, 39], or object-centric and slot decompositions [40]. Discrete VQ-VAEs regularize the latent through a learned codebook [41, 42]; PIWM builds on this quantization but, unlike it, ties codebook components to physical values. These methods improve organization or modularity, yet none guarantees correspondence to measurable physical quantities.

Closer to physical grounding, several models steer latents toward physics. GOKU-net constrains the latent dynamics to a known ODE and infers that ODE’s parameters [43], but the functional form must be fixed a priori and the latents are tied to it rather than to quantities that can be measured from the observation; SPARTAN produces semantically structured outputs but incorporates neither known dynamics nor action-conditioned prediction [44]; Deep Variational Bayes Filters add latent state-space dynamics to VAEs and report latents that correlate with physical quantities [45], but the correspondence emerges from the transition structure rather than being enforced, so which dimension carries which quantity is not fixed in advance; and continuous-time models such as Neural ODEs learn smooth but uninterpretable latent dynamics [46]. Representation learning specialized for pixel-based control has been studied extensively for its effect on downstream performance [47], while again leaving the latent opaque. Compared with these approaches, PIWM explicitly aligns latent dimensions with named physical quantities of both the vehicle and the surrounding road, and constrains their evolution with known physics under only weak, distribution-based supervision.

2.3. World Models

World models learn latent dynamics for imagination-based planning and control. The Dreamer and PlaNet line learns recurrent latent state-space

models from pixels [48, 49, 3]; transformer- and masked-prediction world models improve sample efficiency [2, 4], and the same family has been trained directly on physical robots [50]; and a range of variants explore alternative latent factorizations [1, 51, 52]. These models excel empirically but learn latent spaces with no clear physical meaning, which limits inspection and safety monitoring.

For driving, recent world models predict rich structured futures such as 3D occupancy grids or rendered scenes [53, 54, 55, 56], improving forecasting realism but still without aligning the internal state to physical variables. Neuro-symbolic world models improve generalization, but their symbolic vocabulary is either supplied in advance or induced separately from the physical quantities a controller acts on [57, 58]. The closest prior work, Vid2Param, recovers physical parameters from video [59] but requires full supervision and provides no explicit world model for long-horizon rollout. Compared with these approaches, PIWM recovers physical state under weak supervision and rolls it forward through interpretable dynamics, and, new in this article, it does so for the external road environment, not only the controlled body.

2.4. Physics-Informed and Interpretable Learning

A complementary thread injects physics directly into learning. Sparse identification of nonlinear dynamics and differentiable-physics or end-to-end system identification recover governing equations or parameters [60, 61, 62], while physics-informed neural networks embed differential constraints in the loss [63, 64]. Within world modeling, physical priors have been added as bounds on states and actions [65, 66], as known physics entering the model’s drift term as an inductive bias [67], kinematics-inspired layers [68], and Taylor-monomial-structured latent dynamics (Phy-Taylor) [69].

Most of these methods assume access to the underlying low-dimensional state and are not designed for raw visual inputs, or they enforce physical plausibility without yielding a state that names specific quantities. PIWM differs in two ways aligned with this article’s contributions: it extracts the physical state from images under weak supervision, so that physics can be applied without ground-truth labels, and it applies a structured, learnable physical model, a bicycle model for the vehicle and a geometry-driven transition for the road, to roll the interpretable state forward. Compared with physics-informed predictors that model only the controlled body, PIWM applies interpretable physical structure to the external environment as well.

2.5. Partial Observability, Privileged Learning, and Road Structure

Partial observability is a long-standing challenge in robotics and reinforcement learning, classically formalized through partially observable Markov decision processes and belief-state planning [70, 71]. Model-free agents commonly cope with it by adding memory through recurrent or temporal architectures: recurrent policy gradients [72], gated recurrent networks [73], convolutional LSTMs for spatiotemporal inputs [74], and temporal convolutional networks [75]. These approaches recover an opaque belief or memory state; they do not produce a physically interpretable estimate of the missing environmental variables, which is exactly what our road-context state provides.

When information needed for learning is available only at training time, privileged-information and teacher–student learning supervise a model that must operate without it at deployment [76, 77]. This is the mechanism we use on the real DonkeyCar: ground-truth road context from training-time track localization teaches the encoder what to extract, while the deployed car observes only its front camera and must reconstruct that context from images.

In driving specifically, a substantial literature encodes lane and road structure in learned representations [17, 19], but typically requires explicit maps or HD annotations; and physics-based predictors apply kinematic or dynamic bicycle models to low-dimensional state vectors [78] rather than to learned image representations. Learning-based control has also been demonstrated on RC-scale physical cars, including reinforcement learning on the DonkeyCar platform [79], but without an interpretable predictive world model. Compared with all of the above, our work is, to our knowledge, the first to integrate a factorized vehicle–road representation with a physics-informed world model trained from images under weak supervision, and to both *encode* and *predict* the road environment in physically interpretable terms.

3. Preliminaries

3.1. Autonomous CPS and World Models

Definition 1 (Autonomous CPS). *An autonomous CPS, represented by the tuple $s = (X, I, Y, A, \phi_\theta, g, h)$, models the evolution of an observation-based control system, where X is the physical state space, $I \subset X$ the set of initial states, Y the observation space, A the action space, $\phi_\theta : X \times A \times \Theta \rightarrow X$*

the dynamics with unknown parameters $\theta \in \Theta$, $g : X \rightarrow Y$ the observation function, and $h : Y \rightarrow A$ the controller.

The true state evolves as $x_{t+1} = \phi(x_t, a_t, \theta^*)$, where θ^* are the unknown true parameters. The controller selects actions from observations: $a_t = h(y_t) = h(g(x_t))$. A world model serves as a learned surrogate for forecasting future observations.

Definition 2 (World Model). *A world model $\mathcal{W} = (\mathcal{E}, f, \mathcal{D})$ predicts the future evolution of observations in an autonomous CPS s , where $\mathcal{E} : Y \rightarrow Z$ encodes observations into latent representations, $f : Z \times A \rightarrow Z$ predicts future latents, and $\mathcal{D} : Z \rightarrow Y$ reconstructs predicted observations.*

Conventional world models are evaluated by predictive accuracy in image space (e.g., MSE or SSIM). However, low image error does not imply that the underlying predicted state is physically meaningful: similar-looking frames can correspond to very different physical configurations. For safety monitoring and planning, what matters is not only visual fidelity but whether the latent state can be interpreted as a physical quantity and evolved according to physically plausible dynamics.

3.2. Physical Interpretability

Following the conference paper [8], we define *physical interpretability* via two complementary properties:

1. the latent representation z^* corresponds to meaningful physical variables (state semantics), and
2. its temporal evolution follows physically consistent dynamics (temporal correctness).

Property 1 requires alignment between the latent and the physical state, which PIWM enforces through weak supervision. Property 2 requires that multi-step rollouts of the latent state obey the known equations of motion, which PIWM enforces through the structured dynamics model.

Definition 3 (Interpretable Representation Learning Problem). *Given a dataset $\mathcal{S} = \{(y_{1:M}, a_{1:M}, \Xi_{1:M})\}_{1:N}$, where $\Xi_t = \{\xi_t^{(l)}\}_{l=1}^L$ are samples drawn from an unknown supervisory distribution $p(x_t)$, and a controller h , learn a latent encoder $\mathcal{E}$ such that the mean squared error between the physical state x_t and the encoded latent $z_t^* = \mathcal{E}(y_t)$ is minimized:*

$$\min_{\mathcal{E}} \mathbb{E}[\|x_t - \mathcal{E}(y_t)\|_2^2]. \quad (1)$$

The supervisory samples Ξ_t are drawn from an unknown distribution $p(x_t)$ with no assumed parametric form; we only require that the distribution is *mean-proximal*, meaning its mean lies within a relative interval of width δ around the true value (see Section 6.2 for the precise noise model). This is a strictly weaker assumption than standard supervised learning with precise labels, and represents the uncertainty naturally arising from real-world sensing.

4. Partial Observability in Driving

This section formalizes the challenge introduced in Section 1: in driving, the vehicle state is not a sufficient statistic for the future. The formalization motivates the factorized state that the rest of the article encodes (Section 5.2) and predicts (Section 5.3).

4.1. Why Vehicle State Is Insufficient for Driving

The PIWM conference paper formulated physical interpretability under the implicit assumption that the vehicle’s physical state x_t provides a complete, Markov description of the system: knowing x_t and action a_t determines x_{t+1} . This assumption holds for CartPole, Lunar Lander, and other low-dimensional benchmark tasks, where the system state encodes all task-relevant information. It fails for driving tasks.

In a driving task, the vehicle state $x_t = (x, y, \psi, v)$ (position, heading, speed) does not encode the geometry of the road ahead. The actual lane-following dynamics of the vehicle depend on this geometry:

$$x_{t+1} = \phi(x_t, r_t, a_t, \theta^*), \quad (2)$$

where r_t denotes the *road-context state*, which includes quantities such as the lateral deviation from the lane center, the vehicle’s heading relative to the road tangent, and the local curvature. The road-context component r_t is not directly observed by the onboard camera and is not part of any standard vehicle-state representation.

Definition 4 (Partial Observability in Driving). *A driving task is partially observable with respect to world modeling if the vehicle state x_t alone is not sufficient to determine future motion; equivalently, if there exist two system configurations $(x_t, r_t^{(1)})$ and $(x_t, r_t^{(2)})$ with $r_t^{(1)} \neq r_t^{(2)}$ such that, under the same controller h , the resulting trajectories satisfy $x_{t+1:T}^{(1)} \neq x_{t+1:T}^{(2)}$.*

The road-relative geometry is illustrated in Figure 2. Two vehicles at identical positions on the track, with the same speed and heading, enter different road segments (one straight, one curved). Despite identical vehicle states and identical control actions, the predicted future trajectories diverge, because the curvature of the upcoming road introduces a lateral displacement that the vehicle-state-only dynamics cannot anticipate. Standard PIWM, which learns a single interpretable state aligned with x_t , cannot represent this difference. Two consequences follow, corresponding to the two facets of the challenge from Section 1: the model must *encode* r_t from observations (Section 5.2), and it must *predict* how r_t evolves over the horizon (Section 5.3).

4.2. Factorized Interpretable State

To resolve the partial observability problem, we augment the interpretable state with road-context information and express the vehicle component in a translation- and rotation-invariant reference frame.

Vehicle state. Global coordinates (x, y) introduce unnecessary dependence on the starting position of each trajectory, complicating generalization across scenes. In the road-aligned frame the vehicle’s own motion is described by

$$z_t^v = \begin{bmatrix} v_t \\ \omega_t \end{bmatrix}, \quad (3)$$

its longitudinal speed and yaw rate, both independent of where the trajectory started.

Road-context state. The road-context component z_t^r captures the vehicle’s configuration relative to the road geometry:

$$z_t^r = \begin{bmatrix} s_t \\ e_{\text{CTE}} \\ e_\psi \end{bmatrix}, \quad (4)$$

where s_t is the arc length travelled along the reference centerline, e_{CTE} is the cross-track error (signed lateral distance from the lane center) and e_ψ is the heading error (difference between the vehicle’s heading and the road tangent at its current position). Locating the vehicle by s_t along the road rather than in the world is what allows the road description below to be indexed by it.

Factorized state. The complete *factorized interpretable state* is

$$z_t^* = [z_t^v, z_t^r] = (v_t, \omega_t, s_t, e_{\text{CTE},t}, e_{\psi,t}) \in \mathbb{R}^5. \quad (5)$$

The local curvature κ is deliberately *absent* from (5): it is not propagated. The model instead carries a curvature profile $\hat{\kappa}(\cdot)$ over arc length, perceived once from the image window at the start of the rollout, and reads the curvature at the vehicle from it,

$$\kappa_t = \hat{\kappa}(s_t - s_{t_0}). \quad (6)$$

This distinction is what the rest of the article rests on. A model carrying curvature as a state variable must predict how it evolves, and that prediction compounds over a long rollout; re-indexing a profile observed once cannot compound, because only the vehicle’s position along the road advances.

The Markov property therefore holds for the pair $(z_t^*, \hat{\kappa}(\cdot))$ with $\hat{\kappa}$ fixed across the rollout: given the current factorized state, the perceived profile and the action, the next factorized state is determined up to the unknown parameters θ^* . It does not hold for z_t^* alone, since (6) requires the profile.

4.3. Extended Learning Problem

The learning problem for partially observable driving extends the original PIWM formulation as follows.

Problem 1 (Interpretable State Learning under Partial Observability). *Given a dataset $\mathcal{S} = \{(y_{1:M}, a_{1:M}, \Xi_{1:M}^v, \Xi_{1:M}^r)\}_{1:N}$ where Ξ_t^v and Ξ_t^r are supervisory samples for the vehicle and road-context components respectively, and a controller h , learn a factorized physical encoder $\mathcal{E}_p$ whose two outputs $(z_t^v, z_t^r) = \mathcal{E}_p(y_{1:t})$ approximate the vehicle state and the road-context state, and which supports accurate future prediction under the task dynamics.*

Two aspects distinguish this from the conference formulation, and they map directly onto the two extensions developed next. First, the encoder is conditioned on a *history* $y_{1:t}$ rather than a single frame: road context cannot be inferred from a single image (particularly in the front-camera setting) and requires temporal context (Section 5.2). Second, the supervision is split: Ξ_t^v can often be computed directly from the vehicle state (available in simulation or via external localization), while Ξ_t^r requires knowledge of the road geometry. In the real DonkeyCar experiments, road-context supervision is available only at training time and withheld at test time, so the encoder must learn to recover road context from images alone (Section 5.2), and the predictor must roll it forward (Section 5.3).




Figure 1: PIWM-Frenet computation used in the driving evaluation. (a) The camera history supplies a curvature preview once per rollout; road labels supervise this encoder only during training. (b) The current state, action and indexed curvature drive the structured transition. Predictions are fed back, while the preview remains fixed. The reported rollouts use ground-truth initialization; the figure separates this state input from the camera-derived road preview.

5. PIWM Architecture

The PIWM architecture retains the core structure from the conference version [8], comprising a physically interpretable autoencoder and a learnable dynamics model, and extends it with the factorized state along two axes: a road-context *encoder* (Section 5.2) and a road-context *predictor* (Section 5.3). A high-level overview is shown in Figure 1.

5.1. Learning Interpretable Representations

The autoencoder maps high-dimensional observations to a physically interpretable latent state aligned with the true physical state. We follow the conference paper in considering two architectural approaches.

Intrinsic encoding. The encoder $\mathcal{E} : Y \rightarrow Z^*$ maps an observation directly to the interpretable latent, partitioned as $z^* = [z_p^*, z_{\text{app}}^*]$, where z_p^* is the physically interpretable component and z_{app}^* captures the residual appearance information needed for reconstruction. The conference version writes the latter z_v^* ; we rename it here because this article uses the superscript v for the *vehicle* component of the factorized state. For continuous latent spaces,

the objective follows a β -VAE formulation with interpretability alignment applied to z_p^* :

$$\mathcal{L}_{\text{intrinsic-cont}} = \mathcal{L}_{\text{recon}}(y, \hat{y}) + \lambda_{\text{interp}} \mathcal{L}_{\text{interp}}(z_p^*, \Xi) + \beta D_{\text{KL}}(q(z_{\text{app}}^*|y) \parallel \mathcal{N}(0, I)). \quad (7)$$

For discrete latent spaces, the codebook is partitioned so that each entry $\mathbf{e}_k = [\mathbf{e}_k^p, \mathbf{e}_k^v]$ has a fixed physical component and a learnable visual component. The interpretability loss is applied to the averaged physical part of the selected codebook entries.

Extrinsic encoding. The extrinsic approach decouples visual perception from physical interpretation via a two-stage process. In the first stage, a vision autoencoder $(\mathcal{E}_v, \mathcal{D}_v)$ is trained on image reconstruction only, producing an intermediate latent $z = \mathcal{E}_v(y)$. For continuous variants this is a β -VAE; for discrete variants a VQ-VAE with a 512-entry codebook. In the second stage, the vision encoder is frozen and a physical encoder $\mathcal{E}_p$, paired with a physical decoder $\mathcal{D}_p$ that maps the interpretable state back to the intermediate latent, is trained to produce $z^* = \mathcal{E}_p(z)$, optimizing:

$$\mathcal{L}_{\text{physical}} = \lambda_{\text{interp}} \mathcal{L}_{\text{interp}}(z^*, \Xi) + \lambda_{\text{latent}} \mathcal{L}_{\text{recon}}(z, \mathcal{D}_p(z^*)). \quad (8)$$

The physical encoder $\mathcal{E}_p$ maps the temporal latent sequence to the factorized state $z^* = [z^v, z^r]$ (Section 5.2).

Which encoder for driving. The conference paper’s systematic comparison found that the extrinsic approach with a discrete (VQ) latent gave the strongest physical grounding and prediction accuracy, and it would be natural to carry that configuration over here. We do not, and the reason is worth stating rather than glossing: that comparison was run on *fully observable* benchmarks, where the quantity to be encoded is the controlled body’s own state and a single frame nearly determines it. The road-context encoder faces a different problem—reading the geometry of the scene ahead out of a short, partially informative image window—so the conference ranking is an extrapolation, not a result. We therefore tested it directly, implementing the claimed configuration (per-frame convolutional encoder, a 512-entry VQ bottleneck, and a Transformer with self-attention over the temporal window) at matched capacity and training it to convergence on the same data as the encoder we ship. It is worse at the perception task (RMSE $_{\kappa}$ 0.219 against 0.154 m $^{-1}$) and indistinguishable downstream: paired within folds, its 100-step

position error differs from the shipped encoder’s by less than a millimeter (Table 2). The driving experiments therefore use a convolutional encoder over the stacked window with a linear curvature head, trained from scratch on the road-context objective; the VQ-Transformer variant is reported as a baseline. Section 7.3 returns to why the two differ so much less downstream than they do at perception.

Interpretability loss. For all variants, alignment between the encoded state and the physical variables is enforced through the mean squared error to the empirical mean of the supervisory samples:

$$\mathcal{L}_{\text{interp}}(z^*, \Xi) = \|z^* - \hat{\mu}_{\Xi}\|_2^2, \quad \hat{\mu}_{\Xi} = \frac{1}{L} \sum_{l=1}^L \xi^{(l)}. \quad (9)$$

When applied to the factorized state, this loss is evaluated separately on the vehicle-relative and road-context components with supervision from Ξ^v and Ξ^r :

$$\mathcal{L}_{\text{interp}}^{\text{drive}} = \lambda_v \mathcal{L}_{\text{interp}}(z^v, \Xi^v) + \lambda_r \mathcal{L}_{\text{interp}}(z^r, \Xi^r). \quad (10)$$

The separate weighting of the road-context term λ_r is what forces the encoder to commit capacity to the external environment rather than only to the vehicle.

5.2. Encoding the Road Environment

Challenge and motivation. The conference model encoded only the controlled body’s state, which for the benchmarks was directly inferable from a single frame. Road context is different in two ways. First, it is a property of the *environment*, not the agent, so it must be read off the scene rather than the vehicle’s appearance. Second, for a forward-facing camera it is only *partially visible*: at any instant the camera sees a limited stretch of road, and quantities such as the curvature of the lane the vehicle is about to enter are not fully determined by the current frame. A single-frame encoder is therefore insufficient; the road context must be aggregated over time.

Temporal road-context encoder. We condition the physical encoder on a history of visual latents:

$$(z_t^v, z_t^r) = \mathcal{E}_p(z_{t-k+1:t}), \quad z_{\tau} = \mathcal{E}_v(y_{\tau}), \quad (11)$$

where k is the context-window length. Conditioning on a window rather than a frame is what makes the road context identifiable: as the vehicle




Figure 2: Road state and camera-derived context. (a) The road-relative state is $z^r = (s, e_{CTE}, e_\psi)$: progress along the centerline, signed lateral offset and heading relative to the road tangent. (b) The encoder supplies ten curvature samples over a 4.5m preview. This schematic profile is separate from the propagated state and is held fixed during rollout.

advances, geometry that was previously seen far ahead becomes the geometry the vehicle is on, so the window supplies the parallax and motion cues that disambiguate cross-track error, heading error and curvature, which a single frame leaves ambiguous. In practice we realize $\mathcal{E}_p$ convolutionally, with the k frames of the window entered as input channels and a linear curvature head on the pooled features; a self-attention alternative over the same window was evaluated and did not help (Table 2). The encoder produces both components jointly but supervises them separately through (10), ensuring z_t^r carries genuine environmental information rather than being absorbed into the vehicle component.

Privileged supervision for real-world deployment. On the real DonkeyCar, the road context cannot be measured by the deployed car itself. We therefore use *privileged supervision*: during training, an external localization system (a known track map together with the logged car pose, as is also obtainable from an overhead camera) provides ground-truth road context Ξ^r for each frame; at test time this signal is withheld and the deployed car observes only its front camera. The encoder (11) is thus trained to predict, from the onboard image stream alone, the road context that the privileged signal supplied during training. This is precisely a teacher–student / learning-from-privileged-information setup [76]. Because z^r is physically interpretable, the recovered road context can be inspected and compared against the true lane geometry, which we use as a diagnostic in Section 7.

5.3. Predicting the Road Environment

Challenge and motivation. Encoding the road at the current step is necessary but not sufficient. To roll out a trajectory over a long horizon without new observations, the world model must anticipate how the road context itself changes as the vehicle advances: the curve that is ahead now is underfoot a few steps later, and the cross-track and heading errors evolve as the vehicle tracks (or departs from) the lane. A model that predicts only the vehicle state, holding the road context fixed or ignoring it, drifts quickly because it steers the predicted vehicle through a road that no longer matches reality. We therefore make road prediction an explicit, physically interpretable part of the dynamics.

Split dynamics. We use a *split-dynamics* model that advances the vehicle and the road context with separate, physically appropriate transitions:

$$\hat{z}_{t+1}^v = \phi_{\theta}^v(z_t^v, a_t), \quad (12)$$

$$\hat{z}_{t+1}^r = f_r(z_t^r, z_t^v; \theta_r). \quad (13)$$

The vehicle transition ϕ_{θ}^v is a structured kinematic model with a small learned actuation map (below); the road transition f_r is a geometry-driven model (below). The superscript distinguishes it from the conference version’s ϕ , which advances the whole state from two previous steps; here only the vehicle branch is advanced, and from one. Note which arguments f_r does and does not take. It consumes the vehicle’s *motion*—which over a rollout is the model’s own prediction from the preceding step—and not the action: the road context is advected by how far and how the vehicle moved, so an action influences the road only through the motion it produces. (The one exception is the small bounded residual in (19)–(20), which is conditioned on the action.) This is what makes the second equation a transport rule rather than a second learned dynamics model. Splitting the dynamics decouples the fast-varying vehicle motion from the slow-varying road geometry and lets each receive its correct inductive bias.

Structured kinematics for the vehicle. For driving environments we replace the generic second-order dynamics of the conference paper with a wheeled-vehicle kinematic prior. In the body frame this is the classical *bicycle kinematic*

model

$$x_{t+1} = x_t + v_t \cos(\psi_t + \beta_t) \Delta t, \quad (14)$$

$$y_{t+1} = y_t + v_t \sin(\psi_t + \beta_t) \Delta t, \quad (15)$$

$$\psi_{t+1} = \psi_t + \frac{v_t}{l_r} \sin(\beta_t) \Delta t, \quad (16)$$

with slip angle $\beta_t = \arctan\left(\frac{l_r}{l_f + l_r} \tan(\delta_{f,t})\right)$, where l_f and l_r are the distances from the center of mass to the front and rear axles (unknown and learned during training), $\delta_{f,t}$ is the steering angle, v_t the longitudinal speed and ψ_t the heading.

Because the road context is itself defined relative to the lane (Section 4.2), we express the vehicle in the same road-aligned frame, which yields the equivalent kinematics (18)–(20) below and lets the two branches of the split dynamics share one geometry. In either frame the *pose* kinematics are analytic; what is unknown is how the commanded action drives the vehicle. We therefore learn only a compact actuation map

$$v_{t+1} = v_t + g_\theta^v(z_t^v, z_t^r, a_t), \quad \omega_{t+1} = \omega_t + g_\theta^\omega(z_t^v, z_t^r, a_t), \quad (17)$$

where g_θ^v and g_θ^ω are small multilayer perceptrons mapping throttle to longitudinal acceleration and steering to yaw acceleration. Every remaining quantity in the vehicle branch is a physical variable propagated by an analytic equation.

Structured road-context transition. The road-context state $(e_{\text{CTE}}, e_\psi, \kappa)$ evolves slowly and geometrically. As the vehicle moves by the increment implied by its predicted motion, the road ahead is effectively re-indexed: the lane geometry is advected past the vehicle. We therefore model f_r as a geometry-driven transition that (i) updates the cross-track and heading errors from the vehicle’s predicted lateral and angular motion relative to the road, and (ii) does *not* integrate the curvature forward at all, but re-reads it from the already-perceived preview at the arc length the vehicle has covered.

Concretely, we work in a road-aligned (Frenet) frame in which the vehicle state is $(s, e_{\text{CTE}}, e_\psi, v, \omega)$, with s the arc length along the lane center, v the longitudinal speed and ω the yaw rate. Writing $\kappa_t = \hat{\kappa}(s_t)$ for the perceived curvature at the current arc length and $\eta_t = 1 - e_{\text{CTE},t} \kappa_t$ for the Jacobian of

the road-aligned parameterization, the transition is

$$s_{t+1} = s_t + \frac{v_t \cos e_{\psi,t}}{\eta_t} \Delta t, \quad (18)$$

$$e_{\text{CTE},t+1} = e_{\text{CTE},t} + v_t \sin e_{\psi,t} \Delta t + \varepsilon_{\theta_r}^d, \quad (19)$$

$$e_{\psi,t+1} = e_{\psi,t} + \left(\omega_t - \kappa_t \frac{v_t \cos e_{\psi,t}}{\eta_t} \right) \Delta t + \varepsilon_{\theta_r}^\psi, \quad (20)$$

$$\kappa_{t+1} = \hat{\kappa}(s_{t+1}), \quad (21)$$

where Δt is the sampling interval and $(\varepsilon_{\theta_r}^d, \varepsilon_{\theta_r}^\psi)$ is a small learned residual, tightly bounded, that corrects effects the rigid geometric update misses (e.g., imperfect lane tracking). Equations (18)–(20) are analytic; only the residual and the actuation map (17) are learned.

Equation (21) is the crucial one. The curvature is a *lookup into the curvature profile already perceived at the start of the rollout* (Section 5.2), indexed by how far the predicted vehicle has traveled. It is therefore not a free latent that the dynamics must propagate, and it cannot accumulate error over the horizon: the road is *observed* once and then re-indexed, not extrapolated. This is the structural reason the rollout stays bounded, and it is what distinguishes the model from world models that carry the road as an unconstrained latent that drifts and, in turn, corrupts the vehicle branch that consumes it. The preview is long enough for this to be well defined for most of the horizon: over a 100-step rollout the vehicle covers on average 3.1 m of arc length, staying inside the 4.5 m perceived preview in 82% of evaluation windows; past the preview the model holds the last perceived value rather than extrapolating. Because f_r operates on the physically interpretable z^r , the forecasted road context at any horizon is itself interpretable: the model outputs a predicted lane geometry that can be visualized and checked, not an opaque feature.

5.4. Training Procedure

Training follows a staged pipeline. Stage 1 trains the vision autoencoder $(\mathcal{E}_v, \mathcal{D}_v)$ on image reconstruction. Stage 2 trains the physical encoder $\mathcal{E}_p$ to produce the factorized interpretable state, optimizing (10) jointly with a latent reconstruction term. Stage 3 trains the split-dynamics models $(\phi_\theta^v$ for the vehicle and f_r for the road context), with an optional end-to-end fine-tuning pass that updates the encoder and dynamics together to align the encoded state with what the dynamics can predict. Algorithm 1 summarizes the procedure.




Figure 3: Road prediction by re-indexing. (a) Successive states query the same camera-derived curvature profile at increasing arc-length offsets. (b) Frenet kinematics advance progress, lateral offset and heading error, with the learned corrections described in the text. The profile itself is not extrapolated; perception and state-integration errors can still affect the rollout.

Algorithm 1 PIWM training for partially observable driving.

Input: Training set $\mathcal{S}$, loss weights $\lambda_v, \lambda_r, \lambda_{\text{latent}}$, context length k

- 1: Train vision autoencoder $(\mathcal{E}_v, \mathcal{D}_v)$ on image reconstruction
- 2: **for** each batch $(y_{1:M}, a_{1:M}, \Xi_{1:M}^v, \Xi_{1:M}^r)$ **do**
- 3: $z_t \leftarrow \mathcal{E}_v(y_t)$ for all t
- 4: $(z_t^v, z_t^r) \leftarrow \mathcal{E}_p(z_{t-k+1:t})$ ▷ encode vehicle and road (Sec. 5.2)
- 5: Update $\mathcal{E}_p$ with $\lambda_v \mathcal{L}_{\text{interp}}(z_t^v, \Xi_t^v) + \lambda_r \mathcal{L}_{\text{interp}}(z_t^r, \Xi_t^r) + \lambda_{\text{latent}} \mathcal{L}_{\text{recon}}$
- 6: **end for**
- 7: **for** each batch **do**
- 8: Obtain (z_t^v, z_t^r) from the encoders
- 9: $\hat{z}_{t+1}^v \leftarrow \phi_\theta^v(z_t^v, a_t)$ ▷ bicycle model
- 10: $\hat{z}_{t+1}^r \leftarrow f_r(z_t^r, z_t^v; \theta_r)$ ▷ road prediction (Sec. 5.3)
- 11: Update θ (bicycle) and θ_r (road) on the multi-step rollout loss $\|\hat{z}_{t+1}^v - \hat{\mu}_{\Xi_{t+1}^v}^v\|^2 + \|\hat{z}_{t+1}^r - \hat{\mu}_{\Xi_{t+1}^r}^r\|^2$
- 12: **end for**
- 13: (Optional) End-to-end fine-tune $\mathcal{E}_p, \phi_\theta^v, f_r$ jointly
- 14: **return** trained model parameters

6. Experimental Evaluation

6.1. Environments and Evaluation Tiers

We evaluate across three tiers of increasing difficulty. Tier 1 establishes continuity with the conference paper; Tiers 2 and 3 are the two new partially observable driving case studies (Contribution 1).

Tier 1: Conference benchmarks. The conference evaluation [8] comprises **CartPole**, **Lunar Lander** [80], and the **DonkeyCar simulator**. CartPole and Lunar Lander provide controlled physical systems; the simulator provides a visual driving benchmark with an approximate bicycle prior. Figure 4 shows their original observations. We retain these experiments to document the original representation-learning results, separately from the new road-context evaluation on the physical DonkeyCar.

Tier 2: Partially observable simulated driving (CarRacing). **CarRacing**, from the OpenAI Gym suite [80], is a top-down simulated racing environment in which the vehicle navigates a procedurally generated track. Its observation is a bird’s-eye-view RGB image, 96×96 as the environment returns it, which we resize to 64×64 for the encoder. The vehicle’s kinematic state is available from the simulator, but the road geometry (the shape of the upcoming track and the vehicle’s relationship to the lane) is not part of the state and must be inferred from the image. Because the camera is top-down, the road ahead is visible but the model must still *encode* it into the physical road context and *predict* its evolution during rollout. This is the controlled setting in which the factorized-state formulation and the split dynamics were developed; the quantitative evaluation reported in this article is on the physical platform of Tier 3, for the reason given in Section 7.2.

Tier 3: Real-world DonkeyCar deployment. The **DonkeyCar (real)** tier evaluates the factorized formulation on a physical RC-scale autonomous car driven on a laboratory track. The car observes its environment through a front-facing camera (the test-time observation, 64×64), and logs wheel-encoder speed, steering, and pose. This is the most demanding tier: the camera is forward-facing, so the road ahead is only partially visible, and the road context cannot be measured by the deployed car. We collect driving trajectories on the physical platform and derive ground-truth road context from external track localization, used as privileged supervision at training time only (Section 5.2); at test time the model infers and rolls out road context from the onboard front-camera stream alone.




Figure 4: Original conference environments: CartPole, Lunar Lander and the DonkeyCar simulator. Frames are extracted from the ground-truth rows of the conference qualitative figures [8]; the simulated driving scene is distinct from the physical platform evaluated in Tier 3.

6.2. Experimental Setup

Benchmarks (Tier 1). The data-collection and weak-supervision protocol follows the conference paper. The conference study collected 60,000 trajectories of at least 50 steps per environment, generated by a mix of random actions and a pretrained neural controller. Weak-supervision labels are synthesized by adding biased uniform noise of relative width $\delta \in \{0\%, 5\%, 10\%\}$ to the ground-truth state, as described in Appendix A.

CarRacing (Tier 2). Driving data are generated with a pretrained controller on procedurally generated tracks; each episode records the onboard image, action, and simulator state. Ground-truth labels for the vehicle pose and for the road context are available directly from the simulator and its track geometry. We use exact ground-truth labels for training (corresponding to $\delta = 0$) and evaluate whether the factorized-state formulation improves long-horizon prediction relative to baselines that use only vehicle-state labels.

DonkeyCar real (Tier 3). We drive the physical DonkeyCar on a closed laboratory track and record synchronized front-camera frames (64×64), control actions (steering, throttle), and vehicle pose, at approximately 22 fps, so that $\Delta t = 1/22$ s and the 100-step evaluation horizon corresponds to about 4.5 s of driving. Trajectories are segmented into clean contiguous laps; brief sensing artifacts (dropped or interpolated frames, pose glitches) are filtered out before training. Ground-truth road context is computed from the known track centerline and the logged pose (privileged supervision), and is withheld at test time. The main comparison uses five folds over 71 episodes, totaling

2,483 validation windows; each fold holds out disjoint episodes. Position errors are reported in meters. The legacy 201-window split is used only where explicitly identified.

Architecture and training. The driving experiments use a vision encoder followed by a physical head that produces the factorized state. The head is trained directly against the road-context labels rather than through a reconstruction objective, because road context is available under privileged supervision at training time and nothing has to be learned by reconstructing the image. This is a departure from the conference version’s autoencoding variants, so we implemented those as well and measured them here rather than inheriting their ranking (Table 2); on this task the choice does not move the rollout error. The vision encoder maps each 64×64 observation to a compact latent, and the physical encoder aggregates a context window of $T_{\text{ctx}} = 15$ frames ($\approx 0.7\text{s}$) to infer the road context from image sequences (Section 5.2). Our rollouts are trained on 16-step horizons and evaluated at up to 100 steps, so the reported horizon is more than six times the training horizon; the baselines are trained on 32-step rollouts, i.e. twice ours, which if anything favors them at the horizons where we claim an advantage. Dynamics are trained with Adam, a cosine learning-rate schedule and early stopping on validation loss. All baselines are trained under identical conditions with the same data, the same optimizer and the same evaluation protocol, and with the batch sizes at which each baseline was tuned; we verified that increasing the batch size degrades these small, latency-bound baselines and would therefore have flattered our model.

6.3. Baselines

We compare against the same families used in the conference paper. **DVBF** [45] is a latent state-space model with a locally-linear transition and no physical prior; **GOKU-net** [43] and **Vid2Param** [59] are physics-aware, the first constraining its latent to a known ODE and the second inferring the parameters of a known analytic model; and **SINDYc** [60] is a classical dynamics-discovery baseline operating on the interpretable latent. Every baseline is initialized from the ground-truth vehicle state *and* the ground-truth lane waypoints at the first timestep of the window, so each starts from an explicit description of the road that our model has to recover from pixels. They are then run as latent transition models, consuming only actions over the rollout. One asymmetry among them should be stated plainly, because it

determines their ordering. Vid2Param additionally infers a latent parameter vector from the image window, since inferring physical parameters from video is the whole of that method; it is consequently the only baseline with any visual input, and it is also the strongest one we measure. DVBF and GOKU-net are run without their own observation pathways—DVBF without its filtering step and GOKU-net without its inference of ODE parameters from the observation sequence—so their numbers should be read as characterizing a *dynamics-only* version of each method rather than the method as published. That reduction was made to hold the comparison at the level of the transition model, but it is a limitation of our evaluation, not a property of those methods, and the gap between our model and them is correspondingly easier to obtain than the gap against Vid2Param. To separate the contribution of *perceiving* the road from that of *representing* it, we also evaluate a privileged **oracle** variant of our own model in which κ is read from the known track map instead of the camera; the gap between the two is the cost of perception.

6.4. Metrics

The primary driving metric is the *position error at rollout step k* : the Euclidean distance, in meters, between the predicted and the true vehicle position after k autoregressive steps, averaged over the $N = 201$ held-out validation windows,

$$E_{xy}(k) = \frac{1}{N} \sum_{i=1}^N \|\hat{p}_{i,k} - p_{i,k}\|_2, \quad (22)$$

where $p_{i,k}$ is the ground-plane position of the vehicle at rollout step k of window i , in meters, and $\hat{p}_{i,k}$ the model’s prediction of it. All models are compared in the same planar frame, whatever internal parameterization each uses. We report E_{xy} at $k \in \{25, 50, 100\}$ and plot the full curve $k \mapsto E_{xy}(k)$; all rollouts are initialized from the ground-truth state at $k = 0$ so that the comparison isolates the dynamics rather than the encoder. Because E_{xy} is a distance and not a squared error, its values are directly readable as centimeters of drift. We additionally report the curvature error of the road-context encoder, RMSE_κ in m^{-1} , to evaluate the encoded road directly. For benchmark tasks (Tier 1) we follow the conference protocol and report state RMSE over all state dimensions at a 30-step horizon, together with the relative parameter-recovery error for the environments whose ground-truth physical parameters are known (CartPole, Lunar Lander).

7. Results

7.1. Conference Benchmark Results (Tier 1)

We restore the original conference results [8] in Figures 5–9. These are inherited experiments, not additional runs of PIWM-Frenet. They compare intrinsic and extrinsic encoders with continuous or discrete latents over 30 prediction steps and supervision-noise levels $\delta \in \{0, 5\%, 10\%\}$. The conference protocol uses five-fold cross-validation. State RMSE in these benchmarks is distinct from the meter-valued position error of the physical-car evaluation.

Prediction and architectural choice. Figures 5, 6 and 7 retain all six panels for each environment. Extrinsic PIWM with discrete latents gives the lowest long-horizon error among the extrinsic variants shown. The intrinsic setting has a different ranking: continuous PIWM generally outperforms the discrete variant. Thus the conference results support choosing the latent parameterization together with the encoder architecture, rather than a universal advantage for quantization. The driving encoder comparison in Table 2 revisits this choice for the new road-perception task.

Parameter estimates and visual rollouts. Figure 8 restores the parameter comparison for CartPole and Lunar Lander. The reference lines make both agreement and departures from the known parameters visible as supervision noise changes. We omit the original DonkeyCar wheelbase panel from this recovery comparison because it has no known-parameter reference. Figure 9 retains all five comparison rows from the original visual rollout example, with three display columns selected for legibility. These examples illustrate visual behavior and do not replace the quantitative comparisons above. The original controller diagnostic is also retained in Appendix Appendix B.




Figure 5: Conference CartPole results [8]. State RMSE over 30 steps for (a) extrinsic and (b) intrinsic methods, under three supervision-noise levels. Curves and shaded regions are the original embedded plot images; labels and layout have been reset. Panel-specific vertical limits are retained. The archived figure does not specify the statistical meaning of its shading.




Figure 6: Conference Lunar Lander results [8]. The layout and method groups follow Figure 5. Original curves and shaded regions are preserved; all panels retain the source RMSE range $[0, 2]$.




Figure 7: Conference DonkeyCar-simulator results [8]. The layout follows Figure 5. Original curves and shaded regions are preserved, including the source vertical limit of 2; curves beyond that limit are clipped as in the original. These are 30-step state-RMSE results and must not be read as the 100-step position errors of the physical car.




Figure 8: Conference parameter estimates for the two controlled benchmarks [8]. Each panel retains the original colored intervals, center marks and yellow ground-truth reference. Rows correspond to $\delta = 0, 5\%, 10\%$, from top to bottom. Parameter names and scales follow the source figure; the source does not specify the interval statistic.




Figure 9: Conference visual rollouts at $\delta = 5\%$ [8]. Ground truth and the four original model rows are shown at $k = 5, 15, 30$ for the DonkeyCar simulator and Lunar Lander. Frames are extracted without retouching; the intermediate display columns and the slide footer have been omitted.

7.2. CarRacing (Tier 2)

CarRacing serves in this article as the controlled simulated setting in which the factorized formulation was developed: the road geometry is present in the image but absent from the vehicle state, so the encoder must recover it from pixels even though the top-down view makes the road ahead fully visible. We do not report a quantitative baseline comparison for this tier here. The comparison in an earlier version of this work was run with a vehicle-only dynamics variant rather than the road-context model described in Section 5.3, and re-running it with the present model is left to future work; Section 8 returns to this point. All quantitative claims about the new road-context extension therefore rest on the physical platform of Section 7.3, where the forward-facing camera makes partial observability genuine rather than nominal.

7.3. Real DonkeyCar Results (Tier 3)

We deploy the learned model on the physical DonkeyCar, where the road context is supervised only during training and must be inferred from the front camera at test time. This is the setting the article is about: the camera sees a short stretch of road, the vehicle state alone is non-Markov, and nothing about the track is available to the deployed model beyond what it can read from its own images.

Long-horizon prediction against baselines. Table 1 and Figure 10 report position error over a 100-step (≈ 4.5 s) autoregressive rollout under five-fold cross-validation over the 71 recorded episodes: each fold holds out one fifth of them (301–647 evaluation windows, 2483 in total), every model is trained on that fold’s training episodes only, and we average within a fold before reporting the mean and standard deviation over the five fold-level means. Comparisons *within* a fold are paired—every model rolls out from the same initial states over the same windows—and the paired statistics below are computed that way. Two regimes are visible, and the crossover between them is the central result. For the first two thirds of the horizon the physics-aware baselines are *ahead* of our model: Vid2Param in particular is nearly four times as accurate at step 25 (0.036 m against 0.141 m). This is expected, and it is a consequence of the comparison being generous to them: each baseline is handed the ground-truth lane waypoints at t_0 , so over short horizons it is interpolating from a description of the road that our model had to recover from pixels. Beyond the crossover the ordering inverts and does not invert

Table 1: Position error E_{xy} in meters on the physical DonkeyCar, at three rollout horizons, as mean $\pm$ standard deviation over five-fold cross-validation (lower is better). All models are driven by the same actions and are initialized from the ground-truth state; the baselines additionally receive the ground-truth lane waypoints at t_0 . Of the baselines only Vid2Param also sees the images—DVBF and GOKU-net are run as dynamics-only models (Section 6.3), which is why Vid2Param leads them. Every model, ours included, keeps the epoch with the best $E_{xy}(100)$; the last column also reports how much each model moves when the held-out episodes change. The map-free row recovers position without consulting the track map at all, reading the road geometry from the perception’s shape output instead of indexing a surveyed centerline.

Model	$E_{xy}(25)$	$E_{xy}(50)$	$E_{xy}(100)$
Vid2Param [59]	0.036 $\pm$ 0.008	0.107 $\pm$ 0.008	0.703 $\pm$ 0.142
GOKU-net [43]	0.046 $\pm$ 0.012	0.131 $\pm$ 0.013	0.725 $\pm$ 0.107
DVBF [45]	0.061 $\pm$ 0.011	0.173 $\pm$ 0.031	0.837 $\pm$ 0.131
SINDYc [60]	diverges in all 2483 windows of all five folds		
PIWM-Frenet (ours, perceived κ)	0.141 $\pm$ 0.009	0.245 $\pm$ 0.021	0.463 $\pm$ 0.045
<i>PIWM-Frenet (map-free)</i>	0.146 $\pm$ 0.012	0.258 $\pm$ 0.023	0.516 $\pm$ 0.056

back: on the fold-mean curves our model is ahead of DVBF from step 64 onward, of GOKU-net from step 75 and of Vid2Param from step 82. By step 100 it reaches 0.463 m against 0.703 m for the strongest baseline. Within each fold, where the comparison is paired, our model is ahead at step 100 in *all five* folds against every baseline: by 0.240 m on average against Vid2Param, 0.262 m against GOKU-net and 0.374 m against DVBF. SINDYc diverges in every one of the 2483 evaluation windows of all five folds, so we report it as divergent rather than quoting a value; past the point of divergence the number records only where the rollout was clamped.

The more informative quantity is the spread. Changing which episodes are held out moves our error by ± 0.046 m, against ± 0.192 m for Vid2Param, ± 0.122 m for GOKU-net and ± 0.160 m for DVBF—factors of 2.7 to 4.2. The baselines’ long-horizon error depends substantially on which part of the track they were trained on; ours much less so. We regard this as the more durable of the two claims, because it does not rest on any single trained model.

What the crossover means. The two regimes separate two different things a world model can be good at. Before the crossover the task is essentially smooth extrapolation of the recent motion, and a flexible latent sequence




Figure 10: Physical DonkeyCar prediction under five-fold cross-validation. (a) Mean position-error curves. (b) Error at step 100: small circles show the five fold means, diamonds their average, and horizontal segments their minimum-maximum range. All models use ground-truth initialization. SINDYc is omitted because every evaluated window diverges.

model does it well. After the crossover the rollout has left the stretch of road that the initial observation described, and the question becomes whether the model knows what road it is now on. A model that carries the road as a free latent must propagate it, and any error in that propagation feeds back into the vehicle branch that consumes it; the result is the compounding growth visible in Figure 10. Our model does not propagate the road at all: by (21) it re-reads the curvature that was already perceived at t_0 , indexed by the arc

length covered so far (Figure 3). The vehicle travels 3.1 m on average over the full rollout, within the 4.5 m preview in 82% of the evaluation windows; in the remainder the model holds the last perceived value rather than extrapolating. For most of the horizon the road is therefore *observed*, and error in it cannot accumulate.

The cost of perceiving the road. Because the road-context encoder (Figure 2) is the only part of the model that has to look at pixels, we can price it directly by replacing its output with progressively better sources of κ (Table 2). Feeding the ground-truth curvature profile—perfect perception of exactly the quantity the encoder predicts—yields 0.466 m; the encoder actually used yields 0.463 m, a paired difference of 0.003 m. Reading κ from the known track map instead—privileged information no image-and-action model could have—is *worse*, at 0.505 m; we return to why below. The advantage over the baselines is thus attributable to the *structure* of the representation, not to privileged information or to perception quality.

The table also settles the encoder-architecture question: it does not matter. The ten camera rows span the conference version’s design space—directly supervised regressors from scratch and from the pretrained backbone, the two-stage extrinsic autoencoder, the single-encoder intrinsic variant, the discrete VQ-Transformer configuration the conference paper found strongest on fully observable benchmarks—together with plain LSTM and Transformer sequence models as non-physical benchmarks. Across these, curvature error spans a factor of 7.3 (0.122–0.898 m^{−1}) while the 100-step position error moves by 0.017 m, about 4% and less than any single row’s spread across folds; paired within folds, no camera encoder differs from the shipped one by more than 0.009 m. The ordering does not even carry over: the most accurate curvature estimate (the Transformer) produces the worst rollout of the ten, and the least accurate (the LSTM, 7 times worse) the second best. This is the clearest evidence we have for the article’s central claim. If the rollout depended on continuously estimating the road, an estimate seven times worse would compound over 100 steps. It does not, because the road enters the dynamics as a fixed preview that is re-indexed geometrically (21) rather than integrated: errors in $\hat{\kappa}$ perturb the trajectory but do not accumulate in it.

What does matter is that re-indexing, and the two privileged rows isolate it. Both hold perfect curvature information and differ only in how it is read: the map row looks κ up at the position the rollout has predicted, while the ground-truth profile row reads a preview observed once and indexed by the

Table 2: Where the curvature comes from. All rows use the identical dynamics and differ only in how κ is obtained; five-fold cross-validation, mean $\pm$ standard deviation. RMSE_κ is the encoder’s own curvature error in m^{-1} and $E_{xy}(100)$ the resulting 100-step position error in meters. The first two rows are privileged and not deployable; the remaining ten read the front camera alone and span the conference version’s encoder design space together with its two non-physical sequence benchmarks. Curvature accuracy varies by a factor of 7.3 across the camera rows while the rollout error varies by 4%—less than any single row moves between folds—and the ordering does not carry over: the best curvature estimate gives the worst rollout.

Source of κ	Input	RMSE_κ	$E_{xy}(100)$
Known track map	map + pose	—	0.505 ± 0.043
Ground-truth profile	true κ	—	0.466 ± 0.046
Extrinsic autoencoder	camera	0.193	0.454 ± 0.048
LSTM	camera	0.898	0.456 ± 0.089
Intrinsic autoencoder	camera	0.136	0.460 ± 0.045
CNN, frozen backbone	camera	0.247	0.461 ± 0.053
CNN, warm-started backbone	camera	0.148	0.462 ± 0.052
CNN, from scratch	camera	0.154	0.463 ± 0.045
VQ + Transformer	camera	0.219	0.463 ± 0.051
Transformer	camera	0.122	0.471 ± 0.047

arc length covered. The second is better by 0.039 m, consistently in all five folds—a larger effect than the entire span of the encoder design space above it. Reading a fixed preview beats consulting a perfect map, because a drift in predicted position sends the map lookup to the wrong part of the track while it merely shifts the preview. The corollary is practical: effort spent sharpening road perception has sharply diminishing returns here, and the remaining error is dominated by the vehicle branch.

Robustness to weak supervision. All results above are trained with exact road-context labels. The conference formulation, however, supervises the interpretable state only through noisy distributions, and the physical platform is exactly where such labels are hardest to obtain: they come from an external localization system whose accuracy is finite. We therefore repeat the training under the conference noise model of Appendix A, injecting biased uniform label noise of relative width $\delta \in \{0\%, 5\%, 10\%\}$ (Table 3, Figure 11). The two model families respond in opposite directions. Every baseline degrades as

Table 3: Weak-supervision noise on the physical DonkeyCar. Position error $E_{xy}(100)$ in meters when the labels used for training are corrupted by biased uniform noise of relative width δ (Appendix A), as mean $\pm$ standard deviation over five folds. The structured model improves as the labels degrade; every baseline gets worse. The last row replaces that noise with a zero-mean Gaussian of matched per-dimension magnitude, which reproduces the gain and shows it to be ordinary regularization rather than a property of the noise model.

Model	$\delta = 0\%$	$\delta = 5\%$	$\delta = 10\%$
PIWM-Frenet (ours)	0.463 ± 0.045	0.452 ± 0.035	0.319 ± 0.044
<i>PIWM-Frenet (map-free)</i>	0.516 ± 0.056	0.506 ± 0.049	0.390 ± 0.036
Vid2Param	0.703 ± 0.142	0.919 ± 0.189	0.938 ± 0.211
GOKU-net	0.725 ± 0.107	0.896 ± 0.152	0.864 ± 0.142
DVBF	0.837 ± 0.131	1.067 ± 0.188	1.075 ± 0.177
<i>matched-magnitude Gaussian control, ours at $\delta = 10\%$: 0.313 ± 0.041</i>			

the labels get noisier. Our model improves, monotonically, from 0.463 m at $\delta = 0$ to 0.319 m at $\delta = 10\%$ —a 31% reduction, while Vid2Param worsens by 33% over the same range. The trend holds in all five folds, and our model’s intervals at $\delta = 0$ and $\delta = 10\%$ do not overlap.

This is not a property of the conference formulation’s particular noise. Replacing it with a zero-mean Gaussian of the same per-dimension magnitude, holding everything else fixed, gives 0.313 m against 0.319 m—a paired difference of -0.005 m whose sign is not consistent across folds. Label noise at this magnitude regularizes the structured model, and any zero-mean noise of that magnitude does it. The finding is the *asymmetry* between the model families, not the form of the noise. Consistent with this, the baselines’ selected epochs move sharply earlier as δ grows—several select the fifth epoch at $\delta = 10\%$ —which is what fitting the label noise almost immediately looks like.

Robustness to a corrupted initial state. The study above perturbs the *supervision*; a complementary question is what happens when the state the rollout starts from is itself wrong, as it would be under an imperfect state estimator at deployment. We therefore repeat the evaluation with Gaussian noise of standard deviation σ , in units of each state variable’s own standard deviation, added to the initial state, with ten draws per window (Table 4).




Figure 11: Physical DonkeyCar robustness to corrupted supervision. Each panel shows one training-noise level for PIWM-Frenet; solid curves are means over five folds and shading spans their minimum–maximum range. The same dashed Vid2Param reference, trained with clean labels, is repeated in all panels. The three panels share both axes.

This is a *test-time* perturbation of the state and should not be confused with the δ *training-label* noise of Table 3; the two probe different failure modes. The ranking is preserved at every level: at $\sigma = 1.5$ our model is at 0.823 m against 1.199 m for the strongest baseline, and the absolute gap widens rather than closes. Relative degradation tells a different story, and we report it because it does not favor us: over this range our error grows by $1.78\times$ against Vid2Param’s $1.72\times$. Vid2Param is marginally the less sensitive of the two in relative terms; it is simply working from a larger error throughout.

Table 4: Robustness to a corrupted *initial state*. $E_{xy}(100)$ in meters under Gaussian perturbation of scale σ applied to the state the rollout starts from, mean $\pm$ standard deviation over five folds, ten draws per window. This is a test-time perturbation and is distinct from the training-label noise of Table 3. Baselines use the same symmetrically selected epochs as Table 1, so the $\sigma = 0$ column reproduces it.

Model	$\sigma = 0$	0.25	0.5	1.0	1.5
PIWM-Frenet (ours)	0.462	0.485	0.540	0.685	0.823
Vid2Param	0.699	0.726	0.788	0.977	1.199
GOKU-net	0.724	0.768	0.856	1.097	1.503
DVBF	0.840	0.876	0.974	1.267	1.681

8. Discussion

Bounded rollout comes from not having to predict the road. The horizon at which our model overtakes the baselines (Figure 10) is the horizon at which the initial observation stops describing the road the vehicle is on. The structural answer this article proposes is to make the road not a predicted quantity at all but an *observed* one that is re-indexed by the vehicle’s own predicted motion (21). Vehicle state alone is non-Markov, so the road context must be *encoded*—but once encoded over a preview that outlasts the horizon, it need not be extrapolated, and that is what keeps the error bounded.

Structure, not information. It would be easy to attribute our long-horizon advantage to the privileged supervision used at training time. Table 2 rules this out. Replacing the camera-derived curvature by a lookup into the true track map—the strongest form of privileged information available—makes the 100-step error *worse*, from 0.463 m to 0.505 m, consistently in all five folds, and feeding the ground-truth curvature profile moves it by 0.003 m. The 0.24 m separating our model from the strongest baseline is therefore not privileged information on our side; the baselines were in fact handed ground-truth lane waypoints at t_0 , which our model never sees. The baseline ordering itself points the same way: the single baseline that retains a visual pathway, Vid2Param, is also the best of them, which is the same lesson the road-context encoder teaches. What separates them is that our state variables are road-aligned physical quantities propagated by analytic geometry, rather than latent coordinates propagated by a learned map.

Interpretability of the forecasted environment. Because the road context is physically interpretable, the forecasted road, not just the forecasted vehicle, can be inspected and monitored. The predicted curvature profile can be compared against the true track geometry at every step—the encoder’s own error, $\text{RMSE}_\kappa = 0.154 \text{ m}^{-1}$, is measured in the physical unit of curvature—providing a diagnostic that opaque world models do not offer. This is directly useful for safety monitoring: a forecasted road that departs from plausible geometry is a detectable failure signal, available before the predicted trajectory itself looks wrong. The same property makes the model’s failure modes addressable rather than merely observable. Because every state variable is a physical quantity, the perturbation study of Table 4 can be repeated one variable at a time to ask which one the rollout is actually sensitive to, turning a robustness number into an instruction about where sensing effort belongs. A latent-state model can report that it is sensitive to initialization; it cannot say to *what*.

Privileged training for real deployment. The real DonkeyCar uses privileged road-context supervision at training and withholds it at test. This reflects practical deployability: external localization (a mapped track, or an overhead view) is available in instrumented training environments but not on a freely deployed car. The fact that the learned encoding transfers to the unaided setting is what makes the approach viable on real hardware.

Limitations. First, the road-context variables are specified manually for each environment; a more general approach might discover the decomposition automatically. Second, the vehicle kinematics are a slip-free approximation that ignores tire dynamics, which may limit accuracy at high speeds or on low-friction surfaces; the physical platform used here operates well inside the linear regime, so this assumption is untested at the limit. Third, the argument that the road need not be extrapolated holds only while the horizon stays inside the perceived preview. On our platform the vehicle covers 3.1 m on average per 100-step rollout, but the 4.5 m preview contains the full rollout in only 82% of the evaluation windows; in the rest the model holds the last perceived curvature. A faster vehicle or a longer horizon would leave the preview more often, and the model would then have to either re-perceive or genuinely extrapolate the road; we have not characterized that regime. Fourth, privileged supervision requires an instrumented training environment to obtain ground-truth road context; reducing this requirement

(e.g., via self-supervised road estimation) is open. Fifth, every model’s epoch is selected on the same held-out windows the tables report. We apply that rule symmetrically, which removes an asymmetry present in an earlier version of this evaluation, but symmetric selection is not unbiased selection and the absolute values should be read with that in mind. It also differs from the conference version, which stopped every model early on validation loss. Sixth, the quantitative evaluation reported here is confined to the physical platform: the CarRacing comparison referred to in Section 7.2 was run with a vehicle-only dynamics variant and has not been repeated with the road-context model, so this article does not claim a simulated-domain result. Seventh, two of the three baselines are evaluated in a reduced form. Holding the comparison at the level of the transition model meant running DVBF without its filtering and GOKU-net without its inference of ODE parameters from observations, leaving Vid2Param as the only baseline with a visual pathway. This understates DVBF and GOKU-net, and it is the most likely explanation for their ordering relative to Vid2Param reversing between our simulated and physical experiments; a comparison against full implementations of both would be a stronger test than the one we report. Finally, the method currently assumes a single vehicle and a single lane; multi-agent and multi-lane settings are future work.

Future directions. Natural next steps include scaling PIWM to open-world driving with richer road topology, multiple agents, and additional sensor modalities (LiDAR, radar); treating other vehicles as additional interpretable road-context variables to handle interactions and occlusions; and discovering the road-context structure from data via symbolic regression or causal structure learning, which would reduce the need for environment-specific design.

9. Conclusion

This article extended the Physically Interpretable World Model (PIWM) from fully observable benchmarks to partially observable driving. The central observation is that driving is partially observable: the vehicle’s physical state alone is non-Markov with respect to future trajectory, because the road context (lane geometry and the vehicle’s relationship to it) is absent from any standard state representation and only partially visible to an onboard camera. We addressed this along three axes. We introduced two partially

observable driving case studies, a simulated CarRacing track and a real physical DonkeyCar. We introduced a physically interpretable *encoding* of the road environment, factorizing the interpretable state into a vehicle-relative component and a road-context component inferred from images, including under privileged supervision for real-world deployment. And we introduced a physically interpretable *prediction* of the road environment: a split-dynamics model in which the vehicle advances under structured kinematics while the road is advected past it, re-indexed by the distance the vehicle has covered rather than propagated as a free latent. Experiments on the real DonkeyCar confirm that this treatment of the external environment, not merely of the vehicle, is what keeps long-horizon prediction bounded under partial observability: at 100 steps and under five-fold cross-validation the model reaches 0.463m against 0.703m for the strongest baseline, ahead in every fold and with a fraction of every baseline’s sensitivity to which episodes are held out; under label noise of up to 10% it *improves* while every baseline degrades, and the original benchmark results are retained unchanged. Replacing perceived curvature by a privileged map lookup makes the result worse, and perfect perception moves it by 3mm, so the advantage is attributable to the structure of the representation rather than to the information available to it. Physically interpretable world modeling for driving thus requires aligning latent representations with physical variables *and* decomposing those representations according to the semantic structure of the task, so that the external environment is both encoded and predicted in physically meaningful terms.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this article.

CRedit Authorship Contribution Statement

[CRedit ROLES — TO BE COMPLETED BEFORE SUBMISSION.]

Data Availability

The DonkeyCar trajectories, the derived road-context labels and the code used to train and evaluate the models will be made available in a public repository upon publication.

Acknowledgements

The authors thank Liam (Cade) McGlothlin and Vedansh Maheshwari for their help in exploring interpretable world models. This material is based upon work supported by the U.S. National Science Foundation under Grant Numbers CNS 2513076 and CCF 2403616. Any opinions, findings, and conclusions expressed in this material are those of the authors and do not necessarily reflect the views of the NSF.

Appendix A. Weak Supervision Noise Model

The weak supervision in Tier 1 benchmark experiments follows the noise model from the conference paper. For the i -th state component with valid range size $\mathcal{X}_i$, we generate a randomly shifted center $\tilde{\mu}_i = x_i + \Delta_i$, $\Delta_i \sim \text{Unif}[-\frac{1}{2}\delta|\mathcal{X}_i|, \frac{1}{2}\delta|\mathcal{X}_i|]$, and draw the supervisory distribution as $p_i(x) = \text{Unif}[\tilde{\mu}_i - \frac{1}{2}\delta|\mathcal{X}_i|, \tilde{\mu}_i + \frac{1}{2}\delta|\mathcal{X}_i|]$. At each timestep, $L = 50$ samples are drawn to form the proxy supervision set Ξ_t . The same noise model is applied to the road-context supervision Ξ^r on the physical DonkeyCar in the robustness study of Table 3, where δ scales the valid range of each road-context component. Unless stated otherwise, the Tier 3 results reported in Section 7.3 use exact ground-truth labels at training time ($\delta = 0$), reflecting the availability of training-time track localization; at test time the road context is withheld in every case and inferred from images.

Appendix B. Conference Controller Diagnostic

The conference study [8] passed reconstructed observations through the original fixed controller to assess task-relevant visual information. Table B.5 reproduces the archived results, separately from trajectory-prediction error. The continuous intrinsic variant gives the best controller agreement in these examples; the strongest prediction variant need not have the strongest reconstruction-based control metric.

Table B.5: Conference controller results on reconstructed observations, reproduced from [8]. Entries retain the source’s reported $\pm$ quantities; the archived table does not define their dispersion statistic. Accuracy entries are percentages. These results do not evaluate the new Frenet model.

Architecture	Latent	$\delta = 0$	$\delta = 5\%$	$\delta = 10\%$
<i>DonkeyCar simulator: action RMSE $\downarrow$</i>				
Intrinsic	Continuous	0.12 ± 0.04	0.13 ± 0.04	0.15 ± 0.05
Intrinsic	Discrete	0.21 ± 0.15	0.29 ± 0.16	0.32 ± 0.20
Extrinsic	Continuous	0.15 ± 0.05	0.16 ± 0.05	0.22 ± 0.06
Extrinsic	Discrete	0.15 ± 0.04	0.17 ± 0.05	0.19 ± 0.05
<i>Lunar Lander: action accuracy (%) $\uparrow$</i>				
Intrinsic	Continuous	93.0 ± 1.8	90.5 ± 2.0	87.1 ± 2.2
Intrinsic	Discrete	85.5 ± 2.5	82.1 ± 2.8	78.3 ± 3.1
Extrinsic	Continuous	86.2 ± 2.4	83.5 ± 2.6	80.0 ± 2.9
Extrinsic	Discrete	91.5 ± 2.1	88.6 ± 2.3	84.5 ± 2.5
<i>CartPole: action accuracy (%) $\uparrow$</i>				
Intrinsic	Continuous	98.0 ± 1.0	96.5 ± 1.2	94.0 ± 1.5
Intrinsic	Discrete	95.0 ± 1.6	91.5 ± 2.0	87.2 ± 2.5
Extrinsic	Continuous	95.5 ± 1.5	92.0 ± 1.8	88.0 ± 2.2
Extrinsic	Discrete	97.2 ± 1.1	95.0 ± 1.4	92.5 ± 1.8

References

- [1] D. Ha, J. Schmidhuber, World models, arXiv preprint arXiv:1803.10122 (2018).
- [2] V. Micheli, E. Alonso, F. Fleuret, Transformers are sample-efficient world models, arXiv preprint arXiv:2209.00588 (2022).
- [3] D. Hafner, J. Pasukonis, J. Ba, T. Lillicrap, Mastering diverse domains through world models, arXiv preprint arXiv:2301.04104 (2023).
- [4] Y. Seo, D. Hafner, H. Liu, F. Liu, S. James, K. Lee, P. Abbeel, Masked world models for visual control, in: Conference on Robot Learning, PMLR, 2023, pp. 1332–1344.

- [5] O. Hasan, S. Tahar, Formal verification methods, in: Encyclopedia of Information Science and Technology, Third Edition, IGI global, 2015, pp. 7162–7170.
- [6] S. M. Katz, A. L. Corso, C. A. Strong, M. J. Kochenderfer, Verification of image-based neural network controllers using generative models, *Journal of Aerospace Information Systems* 19 (9) (2022) 574–584.
- [7] M. Alshiekh, R. Bloem, R. Ehlers, B. Könighofer, S. Niekum, U. Topcu, Safe reinforcement learning via shielding, in: Proceedings of the AAAI conference on artificial intelligence, Vol. 32, 2018.
- [8] Z. Mao, M. E. Umasudhan, I. Ruchkin, Physically interpretable world models via weakly supervised representation learning, in: Proceedings of the 17th ACM/IEEE International Conference on Cyber-Physical Systems (ICCPS), Saint Malo, France, 2026, arXiv:2412.12870.
- [9] D. Fridovich-Keil, A. Bajcsy, J. F. Fisac, S. L. Herbert, S. Wang, A. D. Dragan, C. J. Tomlin, Confidence-aware motion prediction for real-time collision avoidance¹, *The International Journal of Robotics Research* 39 (2-3) (2020) 250–265.
- [10] S. Ammoun, F. Nashashibi, Real time trajectory prediction for collision risk estimation between vehicles, in: 2009 IEEE 5Th international conference on intelligent computer communication and processing, IEEE, 2009, pp. 417–422.
- [11] A. Li, L. Sun, W. Zhan, M. Tomizuka, M. Chen, Prediction-based reachability for collision avoidance in autonomous driving, in: 2021 Intl. Conf. on Robotics and Automation (ICRA), 2021.
- [12] K. Nakamura, S. Bansal, Online update of safety assurances using confidence-based predictions, in: 2023 IEEE International Conference on Robotics and Automation (ICRA), IEEE, 2023, pp. 12765–12771.
- [13] A. Muthali, H. Shen, S. Deglurkar, M. H. Lim, R. Roelofs, A. Faust, C. Tomlin, Multi-agent reachability calibration with conformal prediction, in: 2023 62nd IEEE Conference on Decision and Control (CDC), IEEE, 2023, pp. 6596–6603.

- [14] A. Alahi, K. Goel, V. Ramanathan, A. Robicquet, L. Fei-Fei, S. Savarese, Social lstm: Human trajectory prediction in crowded spaces, in: Proc. of the IEEE conference on computer vision and pattern recognition, 2016.
- [15] A. Gupta, J. Johnson, L. Fei-Fei, S. Savarese, A. Alahi, Social gan: Socially acceptable trajectories with generative adversarial networks, in: Proc. of the IEEE conference on computer vision and pattern recognition, 2018.
- [16] S. Wen, H. Wang, D. Metaxas, Social ode: Multi-agent trajectory forecasting with neural ordinary differential equations, in: European Conference on Computer Vision, Springer, 2022, pp. 217–233.
- [17] T. Salzmann, B. Ivanovic, P. Chakravarty, M. Pavone, Trajectron++: Dynamically-feasible trajectory forecasting with heterogeneous data, in: Computer Vision–ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part XVIII 16, Springer, 2020, pp. 683–700.
- [18] G. Revach, N. Shlezinger, X. Ni, A. L. Escoriza, R. J. Van Sloun, Y. C. Eldar, Kalmannet: Neural network aided kalman filtering for partially known dynamics, IEEE Transactions on Signal Processing 70 (2022) 1532–1547.
- [19] M. Itkina, M. Kochenderfer, Interpretable self-aware neural networks for robust trajectory prediction, in: Conference on Robot Learning, PMLR, 2023, pp. 606–617.
- [20] K.-C. Hsu, K. Leung, Y. Chen, J. F. Fisac, M. Pavone, Interpretable trajectory prediction for autonomous vehicles via counterfactual responsibility, in: 2023 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS), IEEE, 2023, pp. 5918–5925.
- [21] L. Lindemann, X. Qin, J. V. Deshmukh, G. J. Pappas, Conformal prediction for stl runtime verification, in: Proceedings of the ACM/IEEE 14th International Conference on Cyber-Physical Systems (with CPS-IoT Week 2023), 2023, pp. 142–153.
- [22] J. Lu, W. Zhan, M. Tomizuka, Y. Hu, Towards generalizable and interpretable motion prediction: A deep variational bayes approach, in:

International Conference on Artificial Intelligence and Statistics, PMLR, 2024, pp. 4717–4725.

- [23] I. Ruchkin, M. Cleaveland, R. Ivanov, P. Lu, T. Carpenter, O. Sokolsky, I. Lee, Confidence composition for monitors of verification assumptions, in: 2022 ACM/IEEE 13th International Conference on Cyber-Physical Systems (ICCPS), IEEE, 2022, pp. 1–12.
- [24] M. Cleaveland, L. Lindemann, R. Ivanov, G. J. Pappas, Risk verification of stochastic systems with neural network controllers, *Artificial Intelligence* 313 (2022) 103782.
- [25] X. Qin, Y. Xia, A. Zutshi, C. Fan, J. V. Deshmukh, Statistical verification of cyber-physical systems using surrogate models and conformal inference, in: 2022 ACM/IEEE 13th International Conference on Cyber-Physical Systems (ICCPS), IEEE, 2022, pp. 116–126.
- [26] M. Zarei, Y. Wang, M. Pajic, Statistical verification of learning-based cyber-physical systems, in: *Proceedings of the 23rd International Conference on Hybrid Systems: Computation and Control*, 2020, pp. 1–7.
- [27] F. Cairolì, L. Bortolussi, N. Paoletti, Neural predictive monitoring under partial observability, in: *Runtime Verification: 21st International Conference, RV 2021, Virtual Event, October 11–14, 2021, Proceedings 21*, Springer, 2021, pp. 121–141.
- [28] M. Waga, E. Castellano, S. Pruekprasert, S. Klikovits, T. Takisaka, I. Hasuo, Dynamic shielding for reinforcement learning in black-box environments, in: *International Symposium on Automated Technology for Verification and Analysis*, Springer, 2022, pp. 25–41.
- [29] A. Stocco, M. Weiss, M. Calzana, P. Tonella, Misbehaviour prediction for autonomous driving systems, in: *Proceedings of the ACM/IEEE 42nd international conference on software engineering*, 2020, pp. 359–371.
- [30] M. Ma, J. Gao, L. Feng, J. Stankovic, Stlnet: Signal temporal logic enforced multivariate recurrent neural networks, *Advances in Neural Information Processing Systems* 33 (2020) 14604–14614.
- [31] Z. Mao, C. Sobolewski, I. Ruchkin, How safe am i given what i see? calibrated prediction of safety chances for image-controlled autonomy,

- in: 6th Annual Learning for Dynamics & Control Conference, PMLR, 2024, pp. 1370–1387.
- [32] Z. Mao, M. E. Umasudhan, I. Ruchkin, How safe will i be given what i saw? calibrated prediction of safety chances for image-controlled autonomy, arXiv preprint arXiv:2508.09346 (2025).
 - [33] G. E. Hinton, S. Osindero, Y. W. Teh, A fast learning algorithm for deep belief nets, *Neural Computation* 18 (2006) 1527–1554.
 - [34] D. P. Kingma, Auto-encoding variational bayes, arXiv preprint arXiv:1312.6114 (2013).
 - [35] I. Higgins, L. Matthey, A. Pal, C. P. Burgess, X. Glorot, M. M. Botvinick, S. Mohamed, A. Lerchner, beta-vae: Learning basic visual concepts with a constrained variational framework., *ICLR (Poster)* 3 (2017).
 - [36] H. Kim, A. Mnih, Disentangling by factorising, in: *International conference on machine learning*, PMLR, 2018, pp. 2649–2658.
 - [37] R. T. Chen, X. Li, R. B. Grosse, D. K. Duvenaud, Isolating sources of disentanglement in variational autoencoders, *Advances in neural information processing systems* 31 (2018).
 - [38] M. Yang, F. Liu, Z. Chen, X. Shen, J. Hao, J. Wang, Causalvae: Disentangled representation learning via neural structural causal models, in: *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2021, pp. 9593–9602.
 - [39] B. Schölkopf, F. Locatello, S. Bauer, N. R. Ke, N. Kalchbrenner, A. Goyal, Y. Bengio, Toward causal representation learning, *Proceedings of the IEEE* 109 (5) (2021) 612–634.
 - [40] M. Mosbach, J. N. Ewertz, A. Villar-Corrales, S. Behnke, Sold: Slot object-centric latent dynamics models for relational manipulation learning from pixels (2025). arXiv:2410.08822.
URL <https://arxiv.org/abs/2410.08822>
 - [41] A. Van Den Oord, O. Vinyals, et al., Neural discrete representation learning, *Advances in neural information processing systems* 30 (2017).

- [42] A. Razavi, A. Van den Oord, O. Vinyals, Generating diverse high-fidelity images with vq-vae-2, *Advances in neural information processing systems* 32 (2019).
- [43] O. Linial, N. Ravid, D. Eytan, U. Shalit, Generative ODE modeling with known unknowns, in: *Proceedings of the Conference on Health, Inference, and Learning, CHIL '21*, Association for Computing Machinery, New York, NY, USA, 2021, pp. 79–94. doi:10.1145/3450439.3451866. URL <https://dl.acm.org/doi/10.1145/3450439.3451866>
- [44] A. Lei, B. Schölkopf, I. Posner, Spartan: A sparse transformer learning local causation, *arXiv preprint arXiv:2411.06890* (2024).
- [45] M. Karl, M. Soelch, J. Bayer, P. van der Smagt, Deep variational Bayes filters: Unsupervised learning of state space models from raw data, in: *International Conference on Learning Representations (ICLR)*, 2017.
- [46] R. T. Chen, Y. Rubanova, J. Bettencourt, D. K. Duvenaud, Neural ordinary differential equations, *Advances in neural information processing systems* 31 (2018).
- [47] M. Tomar, U. A. Mishra, A. Zhang, M. E. Taylor, Learning representations for pixel-based control: What matters and why? (2021). arXiv:2111.07775. URL <https://arxiv.org/abs/2111.07775>
- [48] D. Hafner, T. Lillicrap, I. Fischer, R. Villegas, D. Ha, H. Lee, J. Davidson, Learning latent dynamics for planning from pixels, in: *International conference on machine learning*, PMLR, 2019, pp. 2555–2565.
- [49] D. Hafner, T. Lillicrap, M. Norouzi, J. Ba, Mastering atari with discrete world models, *arXiv preprint arXiv:2010.02193* (2020).
- [50] P. Wu, A. Escontrela, D. Hafner, P. Abbeel, K. Goldberg, Daydreamer: World models for physical robot learning, in: *Conference on robot learning*, PMLR, 2023, pp. 2226–2240.
- [51] L. Zhang, G. Yang, B. C. Stadie, World model as a graph: Learning latent landmarks for planning, in: *International conference on machine learning*, PMLR, 2021, pp. 12611–12620.

- [52] Z. Zhao, M. Qi, H. Ma, Decomposed vector-quantized variational autoencoder for human grasp generation, in: European Conference on Computer Vision, Springer, 2024, pp. 447–463.
- [53] W. Zheng, W. Chen, Y. Huang, B. Zhang, Y. Duan, J. Lu, Occworld: Learning a 3d occupancy world model for autonomous driving, in: European conference on computer vision, Springer, 2024, pp. 55–72.
- [54] C. Min, D. Zhao, L. Xiao, Y. Nie, B. Dai, Uniworld: Autonomous driving pre-training via world models, arXiv preprint arXiv:2308.07234 (2023).
- [55] Z. Yan, W. Dong, Y. Shao, Y. Lu, L. Haiyang, J. Liu, H. Wang, Z. Wang, Y. Wang, F. Remondino, et al., Renderworld: World model with self-supervised 3d label, arXiv preprint arXiv:2409.11356 (2024).
- [56] S. Zuo, W. Zheng, Y. Huang, J. Zhou, J. Lu, Gaussianworld: Gaussian world model for streaming 3d occupancy prediction, arXiv preprint arXiv:2412.10373 (2024).
- [57] J. Balloch, Z. Lin, R. Wright, X. Peng, M. Hussain, A. Srinivas, J. Kim, M. O. Riedl, Neuro-symbolic world models for adapting to open world novelty, arXiv preprint arXiv:2301.06294 (2023).
- [58] Y. Liang, N. Kumar, H. Tang, A. Weller, J. B. Tenenbaum, T. Silver, J. F. Henriques, K. Ellis, Visualpredicator: Learning abstract world models with neuro-symbolic predicates for robot planning, arXiv preprint arXiv:2410.23156 (2024).
- [59] M. Asenov, M. Burke, D. Angelov, T. Davchev, K. Subr, S. Ramamoorthy, Vid2Param: Modeling of dynamics parameters from video, IEEE Robotics and Automation Letters 5 (2) (2020) 414–421. doi:10.1109/LRA.2019.2959476.
- [60] S. L. Brunton, J. L. Proctor, J. N. Kutz, Sparse identification of nonlinear dynamics with control (SINDYc), IFAC-PapersOnLine 49 (18) (2016) 710–715, 10th IFAC Symposium on Nonlinear Control Systems NOLCOS 2016. doi:10.1016/j.ifacol.2016.10.249.
URL <https://www.sciencedirect.com/science/article/pii/S2405896316318298>

- [61] F. de Avila Belbute-Peres, K. Smith, K. Allen, J. Tenenbaum, J. Z. Kolter, End-to-end differentiable physics for learning and control, *Advances in neural information processing systems* 31 (2018).
- [62] D. Yao, C. Muller, F. Locatello, Marrying causal representation learning with dynamical systems for science, *Advances in Neural Information Processing Systems* 37 (2024) 71705–71736.
- [63] S. Cai, Z. Mao, Z. Wang, M. Yin, G. E. Karniadakis, Physics-informed neural networks (pinns) for fluid mechanics: A review, *Acta Mechanica Sinica* 37 (12) (2021) 1727–1738.
- [64] W. Zhong, H. Meidani, Pi-vae: Physics-informed variational auto-encoder for stochastic differential equations, *Computer Methods in Applied Mechanics and Engineering* 403 (2023) 115664. doi:10.1016/j.cma.2022.115664.
URL <https://www.sciencedirect.com/science/article/pii/S0045782522006193>
- [65] R. Tumu, L. Lindemann, T. Nghiem, R. Mangharam, Physics constrained motion prediction with uncertainty quantification, in: *2023 IEEE Intelligent Vehicles Symposium (IV)*, IEEE, 2023, pp. 1–8.
- [66] K. Sridhar, S. Dutta, J. Weimer, I. Lee, Guaranteed conformance of neurosymbolic models to natural constraints, in: *Learning for Dynamics and Control Conference*, PMLR, 2023, pp. 76–89.
- [67] F. Djeumou, C. Neary, U. Topcu, How to learn and generalize from three minutes of data: Physics-constrained and uncertainty-aware neural stochastic differential equations, *arXiv preprint arXiv:2306.06335* (2023).
- [68] H. Cui, T. Nguyen, F.-C. Chou, T.-H. Lin, J. Schneider, D. Bradley, N. Djuric, Deep kinematic models for kinematically feasible vehicle trajectory predictions, in: *2020 IEEE International Conference on Robotics and Automation (ICRA)*, IEEE, 2020, pp. 10563–10569.
- [69] Y. Mao, Y. Gu, L. Sha, H. Shao, Q. Wang, T. Abdelzaher, Phy-Taylor: Partially Physics-Knowledge-Enhanced Deep Neural Networks via NN Editing, *IEEE Transactions on Neural Networks and Learning Systems* 36 (1) (2025) 447–461. doi:10.1109/TNNLS.2023.3325432.
URL <https://ieeexplore.ieee.org/document/10297119>

- [70] L. P. Kaelbling, M. L. Littman, A. R. Cassandra, Planning and acting in partially observable stochastic domains, *Artificial Intelligence* 101 (1–2) (1998) 99–134.
- [71] L. P. Kaelbling, M. L. Littman, A. W. Moore, Reinforcement learning: A survey, *Journal of artificial intelligence research* 4 (1996) 237–285.
- [72] D. Wierstra, A. Förster, J. Peters, J. Schmidhuber, Solving deep memory POMDPs with recurrent policy gradients, in: *International Conference on Artificial Neural Networks (ICANN)*, Springer, 2007, pp. 697–706.
- [73] J. Chung, C. Gulcehre, K. Cho, Y. Bengio, Empirical evaluation of gated recurrent neural networks on sequence modeling, *arXiv preprint arXiv:1412.3555* (2014).
- [74] X. Shi, Z. Chen, H. Wang, D.-Y. Yeung, W.-K. Wong, W.-c. Woo, Convolutional lstm network: A machine learning approach for precipitation nowcasting, *Advances in neural information processing systems* 28 (2015).
- [75] S. Bai, J. Z. Kolter, V. Koltun, An empirical evaluation of generic convolutional and recurrent networks for sequence modeling, *arXiv preprint arXiv:1803.01271* (2018).
- [76] D. Chen, B. Zhou, V. Koltun, P. Krähenbühl, Learning by cheating, in: *Conference on Robot Learning (CoRL)*, PMLR, 2020, pp. 66–75.
- [77] J. Gou, B. Yu, S. J. Maybank, D. Tao, Knowledge distillation: A survey, *International Journal of Computer Vision* 129 (6) (2021) 1789–1819.
- [78] J. Kong, M. Pfeiffer, G. Schilbach, F. Borrelli, Kinematic and dynamic vehicle models for autonomous driving control design, in: *2015 IEEE Intelligent Vehicles Symposium (IV)*, IEEE, 2015, pp. 1094–1099.
- [79] A. Viitala, R. Boney, Y. Zhao, A. Ilin, J. Kannala, Learning to drive (l2d) as a low-cost benchmark for real-world reinforcement learning, in: *2021 20th International Conference on Advanced Robotics (ICAR)*, IEEE, 2021, pp. 275–281.
- [80] G. Brockman, V. Cheung, L. Pettersson, J. Schneider, J. Schulman, J. Tang, W. Zaremba, Openai gym, *arXiv preprint arXiv:1606.01540* (2016).